

Horizontal Scaling

- ❖ Horizontal scaling, also known as scaling out, involves adding more machines or nodes to a system to handle increased load.
- ❖ It distributes the workload across multiple servers to improve performance and availability.
- ❖ This approach enhances capacity by expanding resources rather than upgrading existing hardware.

Key features of horizontal scaling include:

1. **Improved Fault Tolerance:** Distributes the load across multiple machines, reducing the risk of a single point of failure.
2. **Enhanced Performance:** Increases capacity to handle more requests simultaneously by adding more nodes.
3. **Scalability and Flexibility:** Easily adapts to growing workloads by incorporating additional servers or nodes as needed.

To be Install the Metric server K8s

1.Kubectl apply -f <https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml>

```
controlplane $ kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
serviceaccount/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
service/metrics-server created
deployment.apps/metrics-server created
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
```

2. To verify the command:

Kubectl get pods -n kube-system

```
controlplane $ kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-75bdb5b75d-2b6mr	1/1	Running	2 (4m8s ago)	27d
canal-q652m	2/2	Running	2 (4m7s ago)	27d
canal-wzjz6	2/2	Running	2 (4m8s ago)	27d
coredns-5c69dbb7bd-6xvh1	1/1	Running	1 (4m7s ago)	27d
coredns-5c69dbb7bd-xfk71	1/1	Running	1 (4m7s ago)	27d
etcd-controlplane	1/1	Running	2 (4m8s ago)	27d
kube-apiserver-controlplane	1/1	Running	2 (4m8s ago)	27d
kube-controller-manager-controlplane	1/1	Running	2 (4m8s ago)	27d
kube-proxy-dp5fn	1/1	Running	2 (4m8s ago)	27d
kube-proxy-nhmtq	1/1	Running	1 (4m7s ago)	27d
kube-scheduler-controlplane	1/1	Running	2 (4m8s ago)	27d
metrics-server-7ffbc6d68-9bjhx	0/1	Running	0	79s

3. To create the yaml file:

vi depoly.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        resources:
          requests:
            cpu: "0.3"
            memory: "250Mi"
          limits:
            cpu: "0.5"
            memory: "500Mi"
```

4. To create the pods:

Kubectl create -f deploy.yaml

```
controlplane $ kubectl create -f depoly.yaml
deployment.apps/nginx-deployment created
```

5. To check the pods:

Kubectl get pod

```
controlplane $ kubectl get pod
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6c57cf7458-f7fvj  1/1     Running   0           21s
```

6. To verify the resource utilization:

Kubectl describe pod nginx-deployment-6c57cf7458-f7fvj

```
Limits:
  cpu:    500m
  memory: 500Mi
Requests:
  cpu:    300m
  memory: 250Mi
Environment:  <none>
```

7. To Edit the yaml file:

Kubectl edit -n kube-system deployments.apps metrics-server

```
spec:
  containers:
  - args:
    - --kubectl-insecure-tls
    - --kubectl-preferred-address-types=InternalIP
    - --cert-dir=/tmp
    - --secure-port=10250
    - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
    - --kubelet-use-node-status-port
    - --metric-resolution=15s
    image: registry.k8s.io/metrics-server/metrics-server:v0.7.1
    imagePullPolicy: IfNotPresent
    livenessProbe:
```

8. And edit the images:

Kubectl edit -n kube-system deployments.apps metrics-server

```
controlplane $ kubectl edit -n kube-system deployments.apps metrics-server
deployment.apps/metrics-server edited
```

9. To verify the pods on kubsystem:

Kubectl get pods -n kube-system

```
controlplane $ kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-75bdb5b75d-2b6mr	1/1	Running	2 (16m ago)	27d
canal-q652m	2/2	Running	2 (16m ago)	27d
canal-wzjz6	2/2	Running	2 (16m ago)	27d
coredns-5c69dbb7bd-6xvh1	1/1	Running	1 (16m ago)	27d
coredns-5c69dbb7bd-xfk7l	1/1	Running	1 (16m ago)	27d
etcd-controlplane	1/1	Running	2 (16m ago)	27d
kube-apiserver-controlplane	1/1	Running	2 (16m ago)	27d
kube-controller-manager-controlplane	1/1	Running	2 (16m ago)	27d
kube-proxy-dp5fn	1/1	Running	2 (16m ago)	27d
kube-proxy-nhmtq	1/1	Running	1 (16m ago)	27d
kube-scheduler-controlplane	1/1	Running	2 (16m ago)	27d
metrics-server-774ddc6d5c-kv2xq	0/1	CrashLoopBackOff	4 (15s ago)	106s
metrics-server-7ffbc6d68-9bjhx	0/1	Running	0	14m

10. To using the view to node:

Kubectl top node

```
controlplane $ kubectl top node
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
controlplane	140m	14%	1323Mi	70%
node01	37m	3%	872Mi	46%

1. Create the Yaml file:

vi hpa.yaml

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: nginx-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: nginx-deployment
  minReplicas: 1
  maxReplicas: 10
  targetCPUUtilizationPercentage: 50
```

2. And create the pods:

Kubectl create -f hpa.yaml

```
controlplane $ kubectl create -f hpa.yaml
horizontalpodautoscaler.autoscaling/nginx-hpa created
controlplane $
```

3. You have delete the pods:

Kubectl delete deployments.apps nginx-deployment

```
controlplane $ kubectl delete deployments.apps nginx-deployment
deployment.apps "nginx-deployment" deleted
```

4. Create the yaml file:

vi depoly.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        resources:
          requests:
            cpu: "500m"
            memory: "250Mi"
          limits:
            cpu: "750m"
            memory: "500Mi"
```

5. Create the pod:

Kubectl create -f depoly.yaml

```
controlplane $ kubectl create -f depoly.yaml
deployment.apps/nginx-deployment created
```

6. To verify the pods:

Kubectl get pods

```
controlplane $ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-54c8694f64-g77qk	1/1	Running	0	34s

7. To see the details of scaling:

Kubectrl get horizontalpodautoscalers.autoscaling

```
controlplane $ kubectl get horizontalpodautoscalers.autoscaling
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
nginx-hpa	Deployment/nginx-deployment	cpu: 0%/50%	1	10	1	6m31s

8. To check the status:

Kubectrl get hpa

```
controlplane $ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
nginx-hpa	Deployment/nginx-deployment	cpu: 0%/50%	1	10	1	7m6s

9. To be inside the pod:

Kubectrl exec -it nginx-deployment-54c8694f64-g77qk -- bash

```
controlplane $ kubectl exec -it nginx-deployment-54c8694f64-g77qk -- bash
root@nginx-deployment-54c8694f64-g77qk:/# dd if=/dev/zero of=/dev/null &
[1] 34
root@nginx-deployment-54c8694f64-g77qk:/#
root@nginx-deployment-54c8694f64-g77qk:/# dd if=/dev/zero of=/dev/null &
[2] 35
root@nginx-deployment-54c8694f64-g77qk:/# dd if=/dev/zero of=/dev/null &
[3] 36
root@nginx-deployment-54c8694f64-g77qk:/# dd if=/dev/zero of=/dev/null &
[4] 37
root@nginx-deployment-54c8694f64-g77qk:/# dd if=/dev/zero of=/dev/null &
[5] 38
root@nginx-deployment-54c8694f64-g77qk:/# dd if=/dev/zero of=/dev/null &
[6] 39
root@nginx-deployment-54c8694f64-g77qk:/# exit
exit
```

10. To check the status:

Kubectrl get hpa

```
controlplane $ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
nginx-hpa	Deployment/nginx-deployment	cpu: 75%/50%	1	10	3	11m

11. To check the pods:

Kubectl get pod

```
controlplane $ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-54c8694f64-g77qk	1/1	Running	0	6m29s
nginx-deployment-54c8694f64-n64h9	0/1	Pending	0	53s
nginx-deployment-54c8694f64-vh4mc	1/1	Running	0	68s

12. To verify the scaling:

Kubectl get horizontalpodautoscalers.autoscaling -w

Kubectl get pod

```
controlplane $ kubectl get horizontalpodautoscalers.autoscaling -w
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
nginx-hpa	Deployment/nginx-deployment	cpu: 75%/50%	1	10	3	19m

```
^Ccontrolplane $ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-54c8694f64-g77qk	1/1	Running	0	14m
nginx-deployment-54c8694f64-n64h9	0/1	Pending	0	9m6s
nginx-deployment-54c8694f64-vh4mc	1/1	Running	0	9m21s

13. To verify the pods & Nodes:

Kubectl top pods

Kubectl top nodes

```
controlplane $ kubectl top pods
```

NAME	CPU(cores)	MEMORY(bytes)
nginx-deployment-54c8694f64-g77qk	751m	6Mi
nginx-deployment-54c8694f64-vh4mc	0m	5Mi

```
controlplane $ kubectl top nodesa
error: unknown command "nodesa"
See 'kubectl top -h' for help and examples
controlplane $ kubectl top nodes
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
controlplane	104m	10%	1285Mi	68%
node01	778m	77%	921Mi	48%

14. To check the nodes:

Kubectl top nodes

```
controlplane $ kubectl top nodes
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
controlplane	113m	11%	1283Mi	68%
node01	777m	77%	920Mi	48%

15. To check the pod:

Kubectl top pod

```
controlplane $ kubectl top pod
```

NAME	CPU(cores)	MEMORY(bytes)
nginx-deployment-54c8694f64-g77qk	751m	6Mi
nginx-deployment-54c8694f64-vh4mc	0m	5Mi